

**University of Westminster
School of Computer Science and Engineering
5COSC020W Database Systems**

Module Leader: Ms. Yaalini Balathasan

Individual Coursework

Name: W.A.Heshan Peiris

UOW No.: w1999712

IIT No.: 20221171

Table Of Contents

Contents

Table Of Contents	2
1) Objective	3
2) Scenario.....	3
3) Tasks.....	4
3.1) Data Extraction.....	4
3.2) Data Exploration.....	4
(a) Number of data points.	4
(b) Name of attributes.....	5
(c) Type of attributes	7
(d) Number of missing values for each attribute	9
(e) Entry errors for each attribute.....	11
(f) Heatmaps to check missing values.....	12
3.3) Data Transformation	15
Part a.....	15
Part b.....	17
Part c	20
Part d.....	22
3.4 Self-reflection.....	24

1) Objective

Design and implement a data pipeline that extracts data from a data source, performs data transformation, and loads it into a destination storage system. This task simulates a common data engineering scenario.

2) Scenario

You work for a retail company, and your task is to create a data pipeline to process and store customer purchase data or inventory management data. The data comes from multiple CSV files and needs to be transformed before being loaded into a relational database.

3) Tasks

3.1) Data Extraction

I created a small dataset with information about customers' sales transactions and inventory for a retail store, including features like customer ID, customer Name, email in customer table. Product ID, product Name, product Quantity in inventory table. Transaction ID, Card Type, Transaction Date in Sales Transactions table.

3.2) Data Exploration

(a) Number of data points.

Inventory: 1000 Rows

```
In [105]: # (a) Number of data points of Inventory table
num_data_points_in_the_Inventory_table = Inventory.shape[0]
print(f"num_data_points_in_the_Inventory_table: {num_data_points_in_the_Inventory_table}\n")

num_data_points_in_the_Inventory_table: 1000
```

Customers: 1000 Rows

```
In [104]: # (a) Number of data points of Customer table
num_data_points_in_the_Customer_table = Customers.shape[0]
print(f"num_data_points_in_the_Customer_table: {num_data_points_in_the_Customer_table}\n")

num_data_points_in_the_Customer_table: 1000
```

Sales Transactions: 1000 Rows

```
In [106]: # (a) Number of data points of Sales Transactions table
num_data_points_in_the_Sales_Transactions_table = Sales_Transactions.shape[0]
print(f"num_data_points_in_the_Sales_Transactions_table: {num_data_points_in_the_Sales_Transactions_table}\n")

num_data_points_in_the_Sales_Transactions_table: 1000
```

(b) Name of attributes.

Inventory:

Product ID
Product Name
Last Restock Date
Product Quantity
Unit Price
Supplier

```
In [73]: # (b) Name of attributes of Inventory table
attributes_name_Inventory_table = Inventory.columns
print(f"attributes_name_of_Inventory_table: {attributes_name_Inventory_table}\n")

attributes_name_of_Inventory_table: Index(['product_id', 'product_name', 'last_restock_date', 'product_quantity',
      'unit_price', 'supplier'],
      dtype='object')
```

Customers:

Customer ID
Customer Name
Age
Email
Address
Phone Number
Income

```
In [72]: # (b) Name of attributes of Customer table
attributes_name_of_Customer_table = Customers.columns
print(f"attributes_name_of_Customer_table: {attributes_name_of_Customer_table}\n")

attributes_name_of_Customer_table: Index(['customer_id', 'customer_name', 'age', 'email', 'address',
      'phone_number', 'income'],
      dtype='object')
```

Sales Transactions:

Transaction ID

Customer ID

Product ID

Quantity

Transaction Date

Total Amount

Card Type

```
In [74]: # (b) Name of attributes of Sales_Transactions table
attributes_name_Sales_Transactions_table = Sales_Transactions.columns
print(f"attributes_name_of_Sales_Transactions_table: {attributes_name_Sales_Transactions_table}\n")

attributes_name_of_Sales_Transactions_table: Index(['transaction_id', 'customer_id', 'product_id', 'quantity',
            'transaction_date', 'total_amount', 'Card_Type'],
            dtype='object')
```

(c) Type of attributes

Inventory:

Product ID - **Integer**

Product Name - **String**

Last Restock Date - **Date**

Product Quantity - **Integer**

Unit Price - **Float**

Supplier – **String**

```
In [75]: # (c) Type of attributes of Customer table
attribute_types_of_Customer_table = Customers.dtypes
print("attribute_types_of_Customer_table:")
print(attribute_types_of_Customer_table)

attribute_types_of_Customer_table:
customer_id      object
customer_name    object
age              float64
email            object
address          object
phone_number     object
income           float64
dtype: object
```

Customers:

Customer ID - **Integer**

Customer Name - **String**

Age - **Integer**

Email - **String**

Address - **String**

Phone Number - **Integer**

Income – **Float (Currency)**

```
In [76]: # (c) Type of attributes of Inventory table
attribute_types_of_Inventory_table = Inventory.dtypes
print("attribute_types_of_Inventory_table:")
print(attribute_types_of_Inventory_table)

attribute_types_of_Inventory_table:
product_id      object
product_name     object
last_restock_date object
product_quantity float64
unit_price       float64
supplier         object
dtype: object
```

Sales Transactions:

Transaction ID - **Integer**

Customer ID - **Integer**

Product ID - **Integer**

Quantity - **Integer**

Transaction Date - **Date**

Total Amount - **Float (Currency)**

Card Type – **String**

```
In [77]: # (c) Type of attributes of Sales Transactions table
attribute_types_of_Sales_Transactions_table = Sales_Transactions.dtypes
print("attribute_types_of_Sales_Transactions_table:")
print(attribute_types_of_Sales_Transactions_table)

attribute_types_of_Sales_Transactions_table:
transaction_id    object
customer_id       object
product_id        object
quantity          float64
transaction_date   object
total_amount      int64
Card_Type         object
dtype: object
```


(d) Number of missing values for each attribute

Inventory:

Product ID - **0**

Product Name - **0**

Last Restock Date - **109**

Product Quantity - **110**

Unit Price - **102**

Supplier – **0**

```
In [114]: # (d) Number of missing values for each attribute
missing_values_of_Inventory_table = Inventory.isnull().sum()
print("\nNumber of missing values for each attribute:")
print(missing_values_of_Inventory_table)
```

```
Number of missing values for each attribute:
product_id          0
product_name        0
last_restock_date   109
product_quantity    110
unit_price          102
supplier            0
dtype: int64
```

Customers:

Customer ID - **0**

Customer Name - **0**

Age - **109**

Email - **114**

Address - **101**

Phone Number - **0**

Income – **257**

```
In [113]: # (d) Number of missing values for each attribute
missing_values_of_Customer_table = Customers.isnull().sum()
print("\nNumber of missing values for each attribute:")
print(missing_values_of_Customer_table)
```

```
Number of missing values for each attribute:
customer_id      0
customer_name    0
age             109
email           114
address         101
phone_number     0
income          257
dtype: int64
```

Sales Transactions:

Transaction ID - **0**

Customer ID - **0**

Product ID - **114**

Quantity - **112**

Transaction Date - **108**

Total Amount - **0**

Card Type – **276**

```
In [80]: # (d) Number of missing values for each attribute
missing_values_of_Sales_Transactions_table = Sales_Transactions.isnull().sum()
print("\nNumber of missing values for each attribute:")
print(missing_values_of_Sales_Transactions_table)
```

```
Number of missing values for each attribute:
transaction_id    0
customer_id       0
product_id       114
quantity         112
transaction_date  108
total_amount      0
Card_Type        276
dtype: int64
```

(e) Entry errors for each attribute

```
In [116]: # Get the entry errors for each attribute
entry_errors = (Customers.apply(pd.unique).apply(len) - 1).to_list()
entry_errors
```

```
Out[116]: [999, 999, 73, 886, 899, 999, 743]
```

```
In [117]: # Get the entry errors for each attribute
entry_errors = (Inventory.apply(pd.unique).apply(len) - 1).to_list()
entry_errors
```

```
Out[117]: [999, 817, 338, 853, 898, 999]
```

```
In [118]: # Get the entry errors for each attribute
entry_errors = (Sales_Transactions.apply(pd.unique).apply(len) - 1).to_list()
entry_errors
```

```
Out[118]: [999, 999, 886, 847, 330, 998, 16]
```

(f) Heatmaps to check missing values

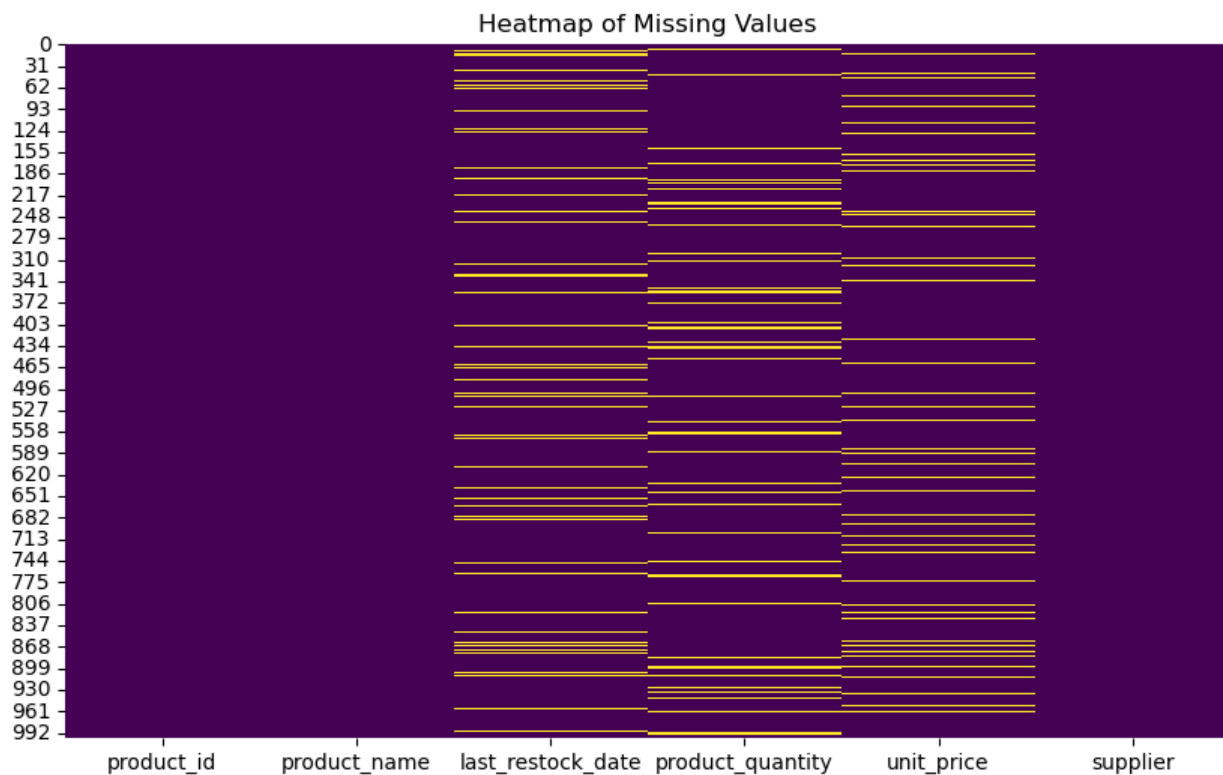
Customer Table

```
In [82]: # (f) Heatmap to check missing values
plt.figure(figsize=(10, 6))
sns.heatmap(Customers.isnull(), cbar=False, cmap='viridis')
plt.title('Heatmap of Missing Values')
plt.show()
```



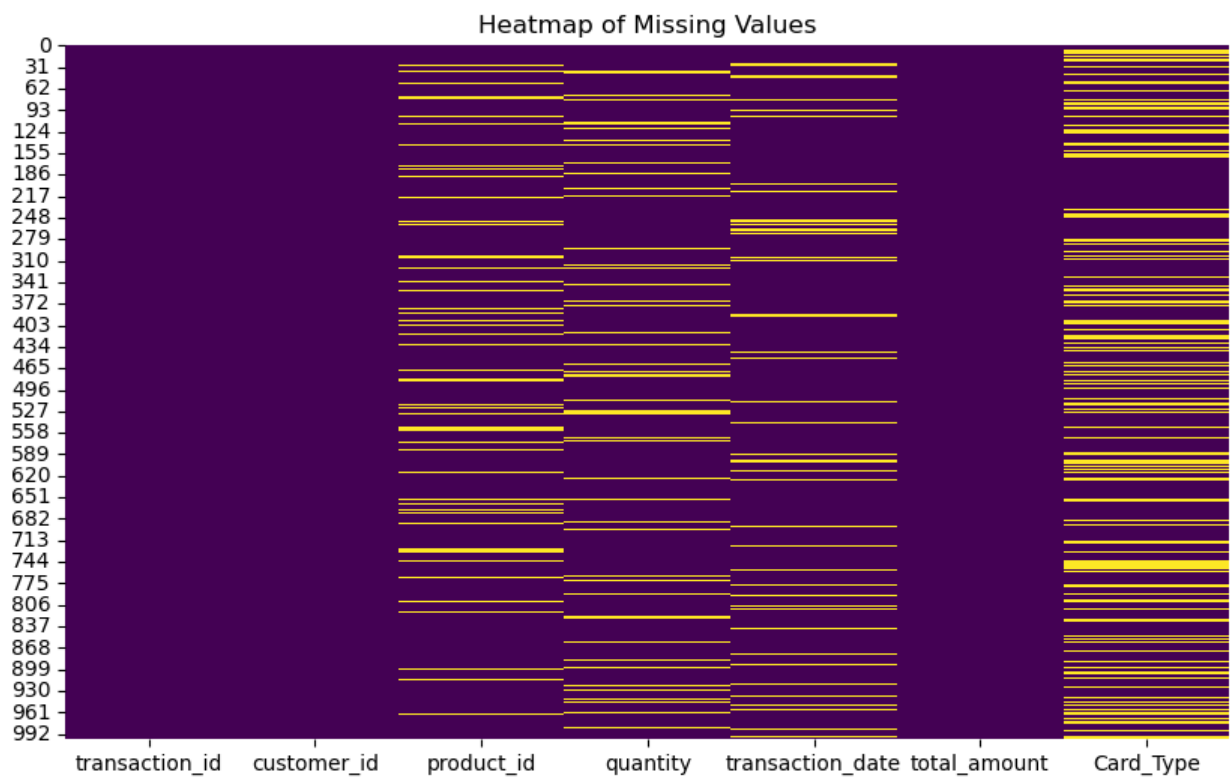
Inventory Table

```
In [83]: # (f) Heatmap to check missing values
plt.figure(figsize=(10, 6))
sns.heatmap(Inventory.isnull(), cbar=False, cmap='viridis')
plt.title('Heatmap of Missing Values')
plt.show()
```



Sales Transactions Table

```
In [85]: # (f) Heatmap to check missing values
plt.figure(figsize=(10, 6))
sns.heatmap(Sales_Transactions.isnull(), cbar=False, cmap='viridis')
plt.title('Heatmap of Missing Values')
plt.show()
```



3.3) Data Transformation

Part a

**Integrating data from multiple CSV Biles of the initial data source into a single dataset.
Eg.combining purchase data from different months, years, etc.**

```
[211]: #3.3(a) Integrating data from multiple CSV Biles of the initial data source into a single dataset  
import pandas as pd  
  
#Reading the 3 data tables from the CSV file  
Customers = pd.read_csv('Customers.csv')  
Inventory = pd.read_csv('Inventory.csv')  
Sales_Transactions = pd.read_csv('Sales_Transactions.csv')
```

```
[212]: #Merging the transactions  
Stocks = Sales_Transactions.merge(Customers, on='customer_id')  
  
# dropping any duplicates  
Stocks = merged_data.select_dtypes(include=['object', 'int64', 'float64']).drop_duplicates()  
  
#Saving the merged data to a new CSV file  
merged_data.to_csv('Stocks.csv', index=False)
```

```
[213]: #Loading the merged data  
Orders = pd.read_csv('Stocks.csv')  
  
#Displaying the merged data  
print(Stocks)
```

	transaction_id	customer_id	product_id_x	quantity	transaction_date	\
0	13-7427019	87-2919006	46-2732354	3952.0	10.06.2023	
1	65-2968490	56-9876147	07-4959214	318.0	15.02.2023	
2	93-7941794	73-2941955	69-8808573	6494.0	27.01.2023	
3	41-5559823	83-9114467	NaN	9650.0	27.05.2023	
4	59-6523445	42-9271721	33-1257533	1821.0	14.01.2023	
..	
995	04-3287363	88-3114549	19-3860523	7349.0	04.09.2023	
996	06-8190016	43-7099027	36-0923368	6155.0	NaN	
997	54-3489274	00-8254353	34-8540914	7498.0	NaN	
998	30-4499724	93-5885871	07-3399157	6729.0	24.11.2022	
999	76-9064451	31-1340840	NaN	4837.0	09.09.2023	

	total_amount	Card_Type	product_id_y	\
0	136314	visa-electron	46-2732354	
1	478949	jcb	07-4959214	
2	217955	jcb	69-8808573	
3	111286	NaN	NaN	
4	166472	jcb	33-1257533	
..	
995	431431	china-unionpay	19-3860523	
996	289473	NaN	36-0923368	
997	47864	NaN	34-8540914	
998	266238	NaN	07-3399157	
999	17636	diners-club-enroute	NaN	

	product_name	last_restock_date	product_quantity	\
0	Onions - Spanish	31.01.2023	5905.0	
1	Scotch - Queen Anne	18.12.2022	2021.0	
2	Wine - Pinot Noir Latour	28.06.2023	7634.0	
3	Arizona - Green Tea	11.09.2023	8355.0	
4	Gingerale - Schweppes, 355 Ml	24.08.2023	6522.0	
..	
995	Foil - Round Foil	12.09.2023	81.0	
996	Sauce - Mint	16.01.2023	7804.0	
997	Chivas Regal - 12 Year Old	05.03.2023	NaN	
998	Lemon Grass	11.03.2023	9226.0	
999	Cheese - Havarti, Roasted Garlic	16.02.2023	NaN	

	unit_price	supplier
0	4220.30	Fleurette Fernier
1	3464.97	Clio Le Grove
2	1955.62	Elmer Bea
3	268.38	Anstice Wycliffe
4	1663.12	Trudy Landy
..
995	2887.52	Maud Harcarse
996	4008.53	Alessandra Farnill
997	1114.83	Beaufort Cusiter
998	3452.85	Alys Prium
999	1597.30	Yance Feak

[1000 rows x 13 columns]

Part b

(b) Handling missing values by either removing rows with missing data or inputting values.

```
In [214]: import pandas as pd

#Reading the merged data
Stocks = pd.read_csv('Stocks.csv')
```

```
In [215]: #Checking the number of rows with missing values
rows_with_null_values = merged_data.isnull().sum()
print(rows_with_null_values)

#Removing rows with missing values
Stocks.dropna(inplace=True)

#Checking the number of rows after removing missing values
print(len(Stocks))
```

```
transaction_id      0
customer_id         0
product_id_x       114
quantity           112
transaction_date    108
total_amount        0
Card_Type           276
product_id_y       114
product_name        0
last_restock_date   109
product_quantity    110
unit_price          102
supplier            0
dtype: int64
365
```

```
In [216]: #Saving the cleaned data to a new csv file
Stocks.to_csv('Stocks_cleaned.csv', index=False)

#Reading the cleaned orders data from csv file
Stocks_cleaned = pd.read_csv('Stocks_cleaned.csv')

#Displaying the new cleaned orders csv data
print(Stocks_cleaned)
```

	transaction_id	customer_id	product_id_x	quantity	transaction_date	\
0	13-7427019	87-2919006	46-2732354	3952.0	10.06.2023	
1	65-2968490	56-9876147	07-4959214	318.0	15.02.2023	
2	93-7941794	73-2941955	69-8808573	6494.0	27.01.2023	
3	59-6523445	42-9271721	33-1257533	1821.0	14.01.2023	
4	98-9652722	75-4953322	56-9365892	6566.0	13.02.2023	
..	
360	19-6566776	48-0186542	99-8744602	7725.0	04.11.2023	
361	34-5432532	75-9927939	40-8196289	5418.0	01.11.2023	
362	24-9838707	39-9332634	11-0825981	728.0	06.06.2023	
363	73-3038659	75-1150384	69-8539974	9314.0	15.01.2023	
364	04-3287363	88-3114549	19-3860523	7349.0	04.09.2023	

	total_amount	Card_Type	product_id_y	\
0	136314	visa-electron	46-2732354	
1	478949	jcb	07-4959214	
2	217955	jcb	69-8808573	
3	166472	jcb	33-1257533	
4	43692	jcb	56-9365892	
..	
360	126922	americanexpress	99-8744602	
361	231479	jcb	40-8196289	
362	289006	mastercard	11-0825981	
363	40067	maestro	69-8539974	
364	431431	china-unionpay	19-3860523	

	product_name	last_restock_date	product_quantity \
0	Onions - Spanish	31.01.2023	5905.0
1	Scotch - Queen Anne	18.12.2022	2021.0
2	Wine - Pinot Noir Latour	28.06.2023	7634.0
3	Gingerale - Schweppes, 355 Ml	24.08.2023	6522.0
4	Hinge W Undercut	14.04.2023	2700.0
..
360	Water - Tonic	29.04.2023	1970.0
361	Rice - Basmati	29.04.2023	7235.0
362	Pop - Club Soda Can	02.05.2023	3961.0
363	Bread - Mini Hamburger Bun	26.04.2023	2226.0
364	Foil - Round Foil	12.09.2023	81.0

	unit_price	supplier
0	4220.30	Fleurette Fernier
1	3464.97	Clio Le Grove
2	1955.62	Elmer Bea
3	1663.12	Trudy Landy
4	3761.92	Rafaelita Fippe
..
360	2147.33	Marketa Maffy
361	713.01	Wilbur Fishpoole
362	3145.23	Genni Persian
363	3486.25	Charity McCarrison
364	2887.52	Maud Harcarse

[365 rows x 13 columns]

Part c

(c) Removing columns of redundant features.

```
In [217]: #3.3(c)
import pandas as pd

#Loading the cleaned CSV file into a pandas DataFrame
Stocks_cleaned = pd.read_csv('Stocks_cleaned.csv')
```

```
In [218]: #Identifying the columns to remove
columns_to_remove = ['Card_Type', 'total_amount']

#Removing the unwanted columns
Stocks_cleaned.drop(columns=columns_to_remove, axis=1, inplace=True)
```

```
In [219]: #Saving the updated DataFrame to a new CSV file
Stocks_cleaned.to_csv('Updated_data.csv', index=False)

#Reading the cleaned Stocks data
Updated_data = pd.read_csv('Updated_data.csv')

#Displaying the new cleaned Stocks csv data
print(Updated_data)
```

	transaction_id	customer_id	product_id_x	quantity	transaction_date	\
0	13-7427019	87-2919006	46-2732354	3952.0	10.06.2023	
1	65-2968490	56-9876147	07-4959214	318.0	15.02.2023	
2	93-7941794	73-2941955	69-8808573	6494.0	27.01.2023	
3	59-6523445	42-9271721	33-1257533	1821.0	14.01.2023	
4	98-9652722	75-4953322	56-9365892	6566.0	13.02.2023	
..	
360	19-6566776	48-0186542	99-8744602	7725.0	04.11.2023	
361	34-5432532	75-9927939	40-8196289	5418.0	01.11.2023	
362	24-9838707	39-9332634	11-0825981	728.0	06.06.2023	
363	73-3038659	75-1150384	69-8539974	9314.0	15.01.2023	
364	04-3287363	88-3114549	19-3860523	7349.0	04.09.2023	

	product_id_y	product_name	last_restock_date	\
0	46-2732354	Onions - Spanish	31.01.2023	
1	07-4959214	Scotch - Queen Anne	18.12.2022	
2	69-8808573	Wine - Pinot Noir Latour	28.06.2023	
3	33-1257533	Gingerale - Schweppes, 355 Ml	24.08.2023	
4	56-9365892	Hinge W Undercut	14.04.2023	
..	
360	99-8744602	Water - Tonic	29.04.2023	
361	40-8196289	Rice - Basmati	29.04.2023	
362	11-0825981	Pop - Club Soda Can	02.05.2023	
363	69-8539974	Bread - Mini Hamburger Bun	26.04.2023	
364	19-3860523	Foil - Round Foil	12.09.2023	

	product_quantity	unit_price	supplier
0	5905.0	4220.30	Fleurette Fernier
1	2021.0	3464.97	Clio Le Grove
2	7634.0	1955.62	Elmer Bea
3	6522.0	1663.12	Trudy Landy
4	2700.0	3761.92	Rafaelita Fippe
..
360	1970.0	2147.33	Marketa Maffy
361	7235.0	713.01	Wilbur Fishpoole
362	3961.0	3145.23	Genni Persian
363	2226.0	3486.25	Charity McCarrison
364	81.0	2887.52	Maud Harcarse

[365 rows x 11 columns]

Part d

(c) Aggregating purchase data, such as calculating total sales per customer.

```
In [220]: #3.3(d) Aggregating purchase data  
#Creating a new column named Total Amount which depicts the Total by multiplying Unit price by Quantity  
  
import pandas as pd  
  
#Reading the cleaned orders data  
Updated_data = pd.read_csv('Updated_data.csv')  
  
#Creating a column named 'Total Expenditure' and calculating the Total  
Updated_data['Total Expenditure'] = Updated_data['product_quantity'] * Updated_data['unit_price']
```

```
In [221]: #Saving the updated dataframe to a new CSV file  
Updated_data.to_csv('Updated_data_with_total_amount.csv', index=False)  
  
#Reading the updated dataframe with total Total Expenditure  
Updated_data_with_total_amount = pd.read_csv('Updated_data_with_total_amount.csv')  
  
#Displaying the updated dataframe with total Total Expenditure  
print(Updated_data_with_total_Total Expenditure)
```

	transaction_id	customer_id	product_id_x	quantity	transaction_date	\
0	13-7427019	87-2919006	46-2732354	3952.0	10.06.2023	
1	65-2968490	56-9876147	07-4959214	318.0	15.02.2023	
2	93-7941794	73-2941955	69-8808573	6494.0	27.01.2023	
3	59-6523445	42-9271721	33-1257533	1821.0	14.01.2023	
4	98-9652722	75-4953322	56-9365892	6566.0	13.02.2023	
..	
360	19-6566776	48-0186542	99-8744602	7725.0	04.11.2023	
361	34-5432532	75-9927939	40-8196289	5418.0	01.11.2023	
362	24-9838707	39-9332634	11-0825981	728.0	06.06.2023	
363	73-3038659	75-1150384	69-8539974	9314.0	15.01.2023	
364	04-3287363	88-3114549	19-3860523	7349.0	04.09.2023	

	product_id_y	product_name	last_restock_date	\
0	46-2732354	Onions - Spanish	31.01.2023	
1	07-4959214	Scotch - Queen Anne	18.12.2022	
2	69-8808573	Wine - Pinot Noir Latour	28.06.2023	
3	33-1257533	Gingerale - Schweppes, 355 Ml	24.08.2023	
4	56-9365892	Hinge W Undercut	14.04.2023	
..	
360	99-8744602	Water - Tonic	29.04.2023	
361	40-8196289	Rice - Basmati	29.04.2023	
362	11-0825981	Pop - Club Soda Can	02.05.2023	
363	69-8539974	Bread - Mini Hamburger Bun	26.04.2023	
364	19-3860523	Foil - Round Foil	12.09.2023	

	product_quantity	unit_price	supplier	Total Expenditure
0	5905.0	4220.30	Fleurette Fernier	24920871.50
1	2021.0	3464.97	Clio Le Grove	7002704.37
2	7634.0	1955.62	Elmer Bea	14929203.08
3	6522.0	1663.12	Trudy Landy	10846868.64
4	2700.0	3761.92	Rafaelita Fippe	10157184.00
..
360	1970.0	2147.33	Marketa Maffy	4230240.10
361	7235.0	713.01	Wilbur Fishpoole	5158627.35
362	3961.0	3145.23	Genni Persian	12458256.03
363	2226.0	3486.25	Charity McCarrison	7760392.50
364	81.0	2887.52	Maud Harcarse	233889.12

[365 rows x 12 columns]

3.4 Self-reflection

Difficulties

One of the main difficulties I encountered when creating my data pipeline was understanding the different data sources and how they fit together. There were many different data sources, each with its own format and structure, and it was difficult to figure out how to integrate them all into a single pipeline. To assure the quality of the data, I also had to create cleaning and filtering processes because the data was usually incorrect or incomplete.

Another challenge I faced was debugging and troubleshooting the pipeline. There were many points where the pipeline failed, and it was often difficult to identify the source of the problem. This required me to carefully examine the data and the code to find the errors.

Learning

Despite the difficulties, building my data pipeline taught me a lot. The value of design and planning was among the most crucial lessons I took away. I took some time to arrange the various steps of the procedure and the equipment I would need before I began constructing the pipeline. This assisted me in avoiding errors and in building a scalable and effective process.

I also learned a lot about data cleaning and transformation. I developed several techniques for cleaning and transforming the data, such as removing duplicates, handling missing values, and normalizing the data. These techniques were essential for ensuring the quality of the data and for preparing it for analysis.

Along with the technical skills I acquired, I also improved my critical thinking and problem-solving abilities. I gained knowledge on how to tackle challenging issues by decomposing them into smaller, more doable tasks. I also gained knowledge on how to evaluate data critically and spot any issues.

Overall, I found the experience of creating my data pipeline to be challenging but rewarding. I learned a lot about data wrangling and data analysis, and I developed valuable skills that will be useful in my future career.